

Analysis for ad target audience

Cindy Gachuhi

21 January 2022

R Week 1 IP

Defining the question

A Kenyan entrepreneur has created an online cryptography course and would want to advertise it on her blog. She currently targets audiences originating from various countries. In the past, she ran ads to advertise a related course on the same blog and collected data in the process. She would now like to employ your services as a Data Science Consultant to help her identify which individuals are most likely to click on her ads.

Loading the dataset

Dataset: <http://bit.ly/IPAdvertisingData>

```
# Let us import the library 'Data Table'
```

```
#
```

```
library("data.table")
```

```
ad_data <- fread('http://bit.ly/IPAdvertisingData')
```

```
# preview the first 6 values
```

```
head(ad_data)
```

```
##      Daily Time Spent on Site Age Area Income Daily Internet Usage
## 1:                68.95  35    61833.90                256.09
## 2:                80.23  31    68441.85                193.77
## 3:                69.47  26    59785.94                236.50
## 4:                74.15  29    54806.18                245.89
## 5:                68.37  35    73889.99                225.58
## 6:                59.99  23    59761.56                226.74
##
##              Ad Topic Line              City Male   Country
## 1:   Cloned 5thgeneration orchestration Wrightburgh 0   Tunisia
## 2:   Monitored national standardization   West Jodi 1     Nauru
## 3:   Organic bottom-line service-desk     Davidton 0 San Marino
## 4: Triple-buffered reciprocal time-frame West Terrifurt 1     Italy
## 5:      Robust logistical utilization     South Manuel 0     Iceland
## 6:   Sharable client-driven software      Jamieberg 1     Norway
##
##      Timestamp Clicked on Ad
## 1: 2016-03-27 00:53:11      0
## 2: 2016-04-04 01:39:02      0
## 3: 2016-03-13 20:35:42      0
## 4: 2016-01-10 02:31:19      0
```

```
## 5: 2016-06-03 03:36:18      0
## 6: 2016-05-19 14:30:17      0
```

Cleaning the dataset

Removing missing values

using the function 'colSums', we can identify the total number of missing values in each column

```
#
colSums(is.na(ad_data))

## Daily Time Spent on Site      Age      Area Income
##              0              0              0
##      Daily Internet Usage      Ad Topic Line      City
##              0              0              0
##              Male      Country      Timestamp
##              0              0              0
##      Clicked on Ad
##              0
```

We can see that no columns in our dataset contain missing values.

Removing duplicates

To get rid of duplicates, we will identify the unique values from our dataset and assign them to variable 'unique_ad'

```
unique_ad <- unique(ad_data)
```

print the variable and check out the unique values

```
unique_ad

##      Daily Time Spent on Site      Age      Area Income      Daily Internet Usage
## 1:      68.95      35      61833.90      256.09
## 2:      80.23      31      68441.85      193.77
## 3:      69.47      26      59785.94      236.50
## 4:      74.15      29      54806.18      245.89
## 5:      68.37      35      73889.99      225.58
## ---
## 996:      72.97      30      71384.57      208.58
## 997:      51.30      45      67782.17      134.42
## 998:      51.63      51      42415.72      120.37
## 999:      55.55      19      41920.79      187.95
## 1000:      45.01      26      29875.80      178.35
##      Ad Topic Line      City      Male
## 1:      Cloned 5thgeneration orchestration      Wrightburgh      0
## 2:      Monitored national standardization      West Jodi      1
## 3:      Organic bottom-line service-desk      Davidton      0
## 4:      Triple-buffered reciprocal time-frame      West Terrifurt      1
## 5:      Robust logistical utilization      South Manuel      0
## ---
## 996:      Fundamental modular algorithm      Duffystad      1
```

```
## 997:      Grass-roots cohesive monitoring      New Darlene      1
## 998:      Expanded intangible solution      South Jessica      1
## 999:      Proactive bandwidth-monitored policy      West Steven      0
## 1000:      Virtual 5thgeneration emulation      Ronniemouth      0
##          Country          Timestamp Clicked on Ad
## 1:      Tunisia 2016-03-27 00:53:11      0
## 2:      Nauru 2016-04-04 01:39:02      0
## 3:      San Marino 2016-03-13 20:35:42      0
## 4:      Italy 2016-01-10 02:31:19      0
## 5:      Iceland 2016-06-03 03:36:18      0
## ---
## 996:      Lebanon 2016-02-11 21:49:00      1
## 997:      Bosnia and Herzegovina 2016-04-22 02:07:01      1
## 998:      Mongolia 2016-02-01 17:24:57      1
## 999:      Guatemala 2016-03-24 02:35:54      0
## 1000:      Brazil 2016-06-03 21:43:21      1
```

```
# let us confirm whether all the duplicates have been removed
```

```
duplicated_ad <- ad_data[duplicated(ad_data),]
duplicated_ad
```

```
## Empty data.table (0 rows and 10 cols): Daily Time Spent on Site, Age, Area
Income, Daily Internet Usage, Ad Topic Line, City...
```


It seems that the dataset has no duplicated values. We originally had 1000 rows and even after getting the unique values, the number of rows remained the same. After printing out variable 'duplicated_ad', we had no output; confirming that indeed our dataset has no duplicated values.

Checking for outliers


An outlier is an observation that is numerically distant from the rest of the data. When reviewing a boxplot, an outlier is defined as a data point that is located outside the fences ("whiskers") of the boxplot.

```
# using boxplots, we will check for outliers in various columns
# column 'Age'
```


```
boxplot(ad_data$Age) # No outliers!
```




```
#  
boxplot(ad_data$`Daily Time Spent on Site`) # No outliers!
```



```
#  
boxplot(ad_data$`Daily Internet Usage`) # No outliers!
```



```
#  
boxplot(ad_data$`Area Income`)
```



Since the outliers are not too far from our minimum value, we will keep them to avoid removing important data points

Univariate analysis

Measures and Central Tendency

to get the descriptive statistics of all the numerical variables, we will use the function 'summary'
summary(ad_data)

```
## Daily Time Spent on Site      Age      Area Income      Daily Internet
Usage
## Min.      :32.60             Min.      :19.00    Min.      :13996    Min.      :104.8
## 1st Qu.:51.36             1st Qu.:29.00    1st Qu.:47032    1st Qu.:138.8
## Median :68.22             Median :35.00    Median :57012    Median :183.1
## Mean   :65.00             Mean   :36.01    Mean   :55000    Mean   :180.0
## 3rd Qu.:78.55             3rd Qu.:42.00    3rd Qu.:65471    3rd Qu.:218.8
## Max.    :91.43             Max.    :61.00    Max.    :79485    Max.    :270.0
## Ad Topic Line      City      Male      Country
## Length:1000      Length:1000      Min.      :0.000    Length:1000
## Class :character  Class :character  1st Qu.:0.000    Class :character
## Mode  :character  Mode  :character  Median :0.000    Mode  :character
##                               Mean   :0.481
##                               3rd Qu.:1.000
##                               Max.    :1.000
## Timestamp          Clicked on Ad
## Min.      :2016-01-01 02:52:10    Min.      :0.0
```

```
## 1st Qu.:2016-02-18 02:55:42 1st Qu.:0.0
## Median :2016-04-07 17:27:29 Median :0.5
## Mean :2016-04-10 10:34:06 Mean :0.5
## 3rd Qu.:2016-05-31 03:18:14 3rd Qu.:1.0
## Max. :2016-07-24 00:22:16 Max. :1.0

#
# since mode doesn't have an inbuilt function, we will create a variable
# 'mode'
#
mode <- function(v) {
  uniqv <- unique(v)
  uniqv[which.max(tabulate(match(v, uniqv)))]
}

# let us get the mode in column 'Age'
age.mode <- mode(ad_data$Age)

#print out mode
age.mode

## [1] 31
```

Measures of Dispersion

```
# we already got the quantiles, minimum and maximum values
# for the rest of the measures, we will focus on column 'daily time spent on
# site'
#
# range
site_time.range <- range(ad_data$`Daily Time Spent on Site`)
site_time.range

## [1] 32.60 91.43

# variance
# The variance is a numerical measure of how the data values is dispersed
# around the mean.
site_time.var <- var(ad_data$`Daily Time Spent on Site`)
site_time.var

## [1] 251.3371

# standard deviation
site_time.sd <- sd(ad_data$`Daily Time Spent on Site`)
site_time.sd

## [1] 15.85361
```

Univariate visualizations

Frequency distribution using a barplot

```
# fetching the 'Age' column
```

```
age <- ad_data$`Age`
```

```
# computing the frequency distribution using the table() function
```

```
age.fr <- table(age)
```

```
age.fr
```

```
## age
```

```
## 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43  
44
```


```
## 6 6 6 13 19 21 27 37 33 48 48 39 60 38 43 39 39 50 36 37 30 36 32 26 23  
21
```

```
## 45 46 47 48 49 50 51 52 53 54 55 56 57 58 59 60 61
```

```
## 30 18 13 16 18 20 12 15 10 9 7 2 6 4 2 4 1
```

```
# using a barplot to visualize this
```

```
barplot(age.fr)
```



```
# From ages 28-37(the more frequent ages), age 31 is the most frequent
```

```
# Age 61 is the Least frequent
```

Bivariate analysis

```
# Correlation Coefficient
```


It is a normalized measurement of how the two are linearly related. If the correlation coefficient is close to 1, it would indicate that the variables are positively linearly related. For -1, it indicates that the variables are negatively linearly related and the scatter plot almost falls along a straight line with negative slope. And for zero, it would indicate a weak linear relationship between the variables.

get the columns

```
daily_usage <- ad_data$`Daily Internet Usage`  
clicked <- ad_data$`Clicked on Ad`  
site_time <- ad_data$`Daily Time Spent on Site`
```

get the correlation coefficients between the variables using cor()
`cor(daily_usage , site_time)`

```
## [1] 0.5186585
```

*# 0.5; indicates a positive correlation although it is a bit weak.
However, we can see that people who have a higher daily internet usage
spend more time on the site*

```
cor(clicked, daily_usage)
```

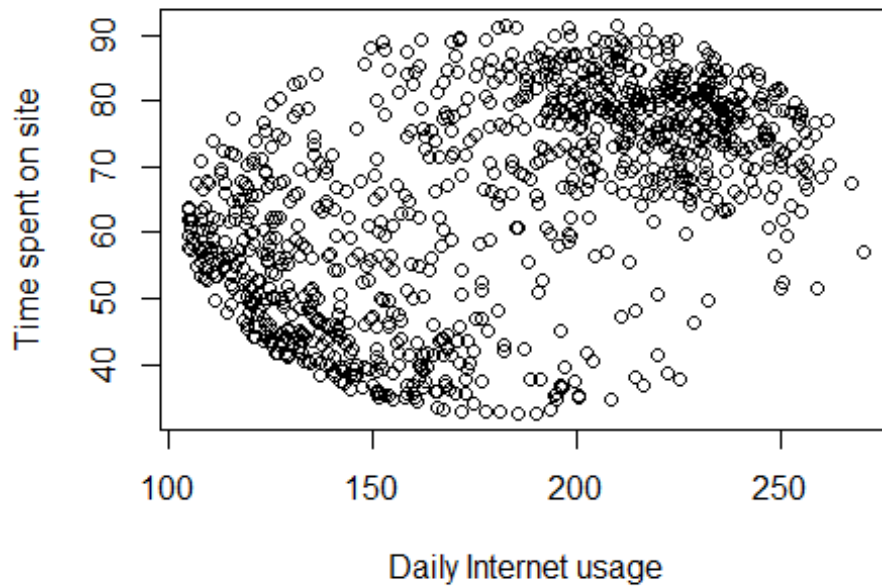
```
## [1] -0.7865392
```

*# -0.78; a negative correlation which is relatively strong.
We see that people who spend more time on the site and on the Internet tend
to click less on the ads.*

Visualizations

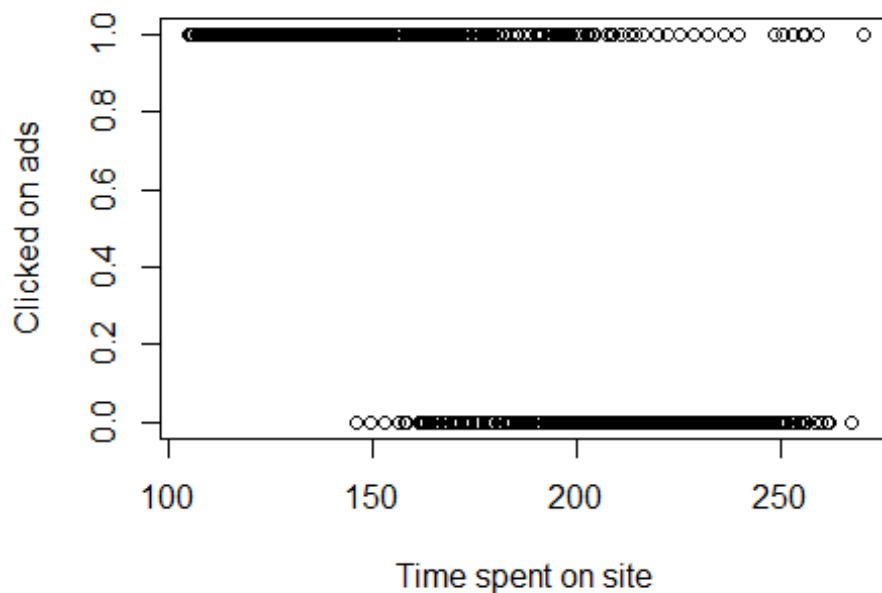
scatterplots

```
plot(daily_usage, site_time, xlab="Daily Internet usage", ylab="Time spent on  
site")
```



*# The points are quite scattered but if you look closely, you see the points show a weak positive linear relationship
between the two variables*

```
plot(daily_usage, clicked, xlab="Time spent on site", ylab="Clicked on ads")
```



With how the points are distributed, we realize a rather strong negative correlation between the variables

From our bivariate analysis, we see that the users who have higher daily internet usage tend to spend more time on the site. Also, we realised that those users who spend more time on the site, actually tend to click on the ads less. This might be because the ads are not in their interest hence they tend to ignore them. A very brief survey should be taken by the user before using the site so as to tailor the ads to the user's interests which might make them click on them more frequently.